

6.828 2014 L1: O/S overview

=

Overview

- * 6.828 goals
 - * Understand operating systems in detail by designing and implementing a small O/S
 - * Hands-on experience with building systems ("Applying 6.033")
- * What do applications want from an O/S?
 - * Abstract the hardware for convenience and portability
 - * Multiplex the hardware among multiple applications
 - * Isolate applications to contain bugs
 - * Allow sharing among applications
- * What is an OS?
 - * e.g. OSX, Windows, Linux
 - * the small view: a h/w management library
 - * the big view: physical machine -> abstract one w/ better properties
- * Organization: layered picture
 - h/w: CPU, mem, disk
 - kernel: [various services]
 - user: applications, e.g. vi and gcc
 - * we care a lot about the interfaces and internal kernel structure
- * What services does an O/S kernel typically provide?
 - * processes
 - * memory
 - * file contents
 - * directories and file names
 - * security
 - * many others: users, IPC, network, time, terminals
- * What does an O/S abstraction look like?
 - * Applications only see them via system calls
 - * Examples, from UNIX / Linux:


```

          ...
          fd = open("out", 1);
          write(fd, "hello\n", 6);
          pid = fork();
          ...
          
```
- * Why is O/S design/implementation hard/interesting?
 - * the environment is unforgiving: weird h/w, no debugger
 - * it must be efficient (thus low-level?)
 - ...but abstract/portable (thus high-level?)
 - * powerful (thus many features?)
 - ...but simple (thus a few composable building blocks?)
 - * features interact: `fd = open(); ...; fork()`
 - * behaviors interact: CPU priority vs memory allocator.
 - * open problems: security, multi-core
- * You'll be glad you learned about operating systems if you...
 - * want to work on the above problems
 - * care about what's going on under the hood
 - * have to build high-performance systems
 - * need to diagnose bugs or security problems

Class structure

- * See web site: <http://pdos.lcs.mit.edu/6.828>

- * Lectures
 - * basic O/S ideas
 - * extended inspection of xv6, a traditional O/S
 - * several more recent topics
 - * xv6 programming to re-inforce xv6 understanding
- * Lab: JOS, a small O/S for x86 in an exokernel style
 - * you build it, 5 labs + final lab of your choice
 - * kernel interface: expose hardware, but protect -- no abstractions!
 - * unprivileged library: fork, exec, pipe, ...
 - * applications: file system, shell, ..
 - * development environment: gcc, qemu
 - * lab 1 is out
 - * make grade
- * Code review
- * Two quizzes: one in class hours, one in final's week

Shell and system calls

- * 6.828 is largely about design and implementation of system call interface. let's start by looking at how programs use that interface. example: the Unix shell.
- * the shell is the Unix command UI
 - * typically handles login session, runs other processes
 - * you saw it in 6.033: <http://web.mit.edu/6.033/www/assignments/hands-on-unix.html>
 - * the shell is also a programming/scripting language
 - * look at some simple examples of shell operations, how they use different O/S abstractions, and how those abstractions fit together. See [Unix paper](../readings/ritchie78unix.pdf) if you are unfamiliar with the shell.
- * [Simplified xv6 sh.c](../homework/sh.c)
 - * See [chapter 0 of xv6 book](../xv6/book-rev8.pdf)
 - * Basic organization: parsing and executing commands (e.g., `ls`, `ls | wc`, `ls > out`)
 - * Shell implemented using system calls (e.g., `read`, `write`, `fork`, `exec`, `wait`)
 - conventions: `-1` return value signals error,
 - error code stored in `<code>errno</code>,`
 - `<code>perror</code>` prints out a descriptive error message based on `<code>errno</code>.`
 - * Many systems calls are encapsulated in libc calls (e.g., `fgets` vs `read`)
- <!--
 Demo:
 - open sh.c in emacs
 - look at `main()`
 - look at `runcmd()`
 - look at `fgets()`
 - man 3 `fgets()`
 -->
- * Trace system calls `$ ls`
 - * On OSX: `sudo dtruss ./a.out` (where `a.out` is the compiled `sh.c`)
 - * On Linux: `strace ./a.out`
- <!--
 - compile sh.c
 - run ./a.out
 - strace ./a.out
 -->
- * what does `fork()` do?
 - copies user memory

- copies process kernel state (e.g. user id)
- child gets a different PID
- child state contains parent PID
- returns twice, with different values
- * parent and child may run concurrently (e.g., on different processors).
- * what does wait() do?
 - waits for any child to exit
 - what if child exits before parent calls wait?

```
<!--  
- strace /bin/sh  
- study output:  
read()  
write()  
stat()  
etc.  
-->
```

- * what are file descriptors? (0, 1, 2, etc. in read/write)

[echo.c](l-overview/echo.c)

- * what is i/o redirection?

[echo.c](l-overview/redirect.c)

How would you implement ">" in sh.c

- * what are pipes? (ls | wc)

[pipe1.c](l-overview/pipe1.c)

[pipe2.c](l-overview/pipe2.c)

How would you implement them in sh.c?

- * Homework assignment for [shell](../homework/xv6-shell.html)