

6.828 2014 L4: Shell & OS organization

--

Homework solution

==

- * Let's review solution for [sh.c](../../homework/sh.c)
 - * exec
 - * redirect
 - * pipe
- * The process graph for a complicated pipeline
 - * Who waits for whom? (draw a tree of processes)
 - * Why close read-end and write-end? ensure that every process starts with 3 file descriptors, and that reading from the pipe returns end of file after the first command exits.

Exploring system calls with more shell features

==

- * You can run the shell, redirect its stdin/stdout, etc.
 - * I'll run this shell script with `sh < script`:


```

          echo one
          echo two
          
```
 - * What will this shell command do?


```

          $ sh < script > out
          
```

 - * the script itself didn't redirect the echo output, but it did inherit a fd 1 that was redirected to out.
 - * to make this work it is important that offset is implicit (maintained by kernel)
 - * this is why read and write don't take an offset argument
 - * Is the following the same as above?


```

          $ echo one > out
          $ echo two > out
          
```
 - * How to implement sequencing/lists (`;&&||`)


```

          $ gcc x.c ; ./a.out
          $ gcc x.c && ./a.out
          
```
 - * How to implement nesting?


```

          $ (echo one; echo two) > out
          
```
 - * How does the shell implement "&"?


```

          $ sleep 2 &
          
```

 - * Q: What if a background process exits while sh waits for a foreground process?

System call observations

==

- * The fork/exec split looks wasteful; why is it useful? (A: exercise 2)

- * System call interface simple, just ints and char buffers. why not have open() return a pointer reference to a kernel file object?
- * Linux has a nice representation of a process and its FDs, under /proc/PID/
 - * maps: VA range, perms (p=private, s=shared), offset, dev, inode, pathname
 - * fd: symlinks to files pointed to by each fd.
- * The file descriptor design:
 - * nice interaction with fork
 - * FDs help make programs more general purpose: don't need special cases for files vs console vs pipe
 - * shell pipelines only work for programs w/ common formats (lines of text)

OS organization

==

- * Now we have a feel for what Unix system call interface provides, how to implement the interface?
- * Why have an o/s at all? why not just a library?
 - then apps are free to use it, or not -- flexible
 - apps can directly interact with hardware
 - some tiny O/Ss for embedded processors work this way
- * Key requirement:
 - isolation
 - multiplexing
 - interaction
- * helpful approach: abstract machine resources rather than raw hardware
 - File system, not raw disk
 - TCP, not a raw ethernet
 - Processes, not raw CPU/memory
 - abstractions often ease multiplexing and interaction
 - also more convenient and portable
- * Goals:
 - apps must use OS interface, cannot directly interact with hardware
 - apps cannot harm operating system
- * Hardware support for isolation
 - * Processors support user mode and kernel mode
 - some instructions can only be executed in kernel mode
 - e.g., instructions to directly interact with hardware
 - * If an application executes a privileged instruction hardware doesn't allow it
 - instead switches to kernel mode
 - kernel can clean up
- * Leverage hardware support
 - * Operating systems runs in kernel mode
 - kernel is a big program
 - services: processes, file system, net
 - low-level: devices, virtual memory
 - all of kernel runs with full hardware privilege (convenient)
 - * Applications run in user mode
 - isolated from kernel
 - systems calls switch between user and kernel mode
 - application call instructions to enter kernel
 - instruction enters kernel at an entry point specified by kernel
- * What to run in kernel mode?

- * xv6 follows a traditional design: all of the OS runs in kernel mode
 - this design is called a monolithic kernel
 - kernel interface ~= system call interface
 - good: easy for subsystems to cooperate
 - one cache shared by file system and virtual memory
 - bad: interactions are complex
 - leads to bugs
 - no isolation within kernel
- * alternative: microkernel design
 - many OS services run as ordinary user programs
 - file system in a file server
 - kernel implements minimal mechanism to run services in user space
 - IPC
 - virtual memory
 - threads
 - kernel interface != system call interface
 - good: more isolation
- * jos: doesn't abstract hardware resources, but still provide isolation

* Can one have process isolation WITHOUT h/w-supported kernel/user mode?
yes!
see Singularity O/S, later in semester
but h/w user/kernel mode is the most popular plan

xv6 kernel address space
==

- * boundary between apps and kernel
 - kernel/user bit
 - kernel address space
 - ...
- * start w. kernel address space
 - machine has booted
 - code runs without virtual memory
 - set up up kernel address space
 - walk through kvmalloc()