

6.828 2014 Lecture 10: Processes, threads, and scheduling

=

Plan:

- homework
- process
- threads
- scheduling

Homework

iderw():

- what goes wrong with adding sti/cli in iderw?
- what ensures atomicity between processors
- what ensures atomicity within a single processor?

filealloc():

- Q: could the disk interrupt handler run while interrupts are enabled?
- Q: does any interrupt handler grab the ftable.lock?
- Q: what interrupt could cause trouble?
- Q: timer interrupt?
- maybe we will find out in this lecture

Process

Process

motivated by isolation

idea: an abstract virtual machine provides the illusion to application of a dedicated computer but an abstract one convenient for application developer one process cannot effect another accidentally

Process API:

- fork
- exec
- exit
- wait
- kill
- sbrk
- getpid

Challenge: more processes than processors

xv6 picture:

- 1 user thread and 1 kernel thread per process
- 1 scheduler thread per processor
- n processors

Terms

- a process: address space plus one or more threads
- a thread: thread of execution
- kernel thread: thread running in kernel mode
- user thread: thread running in user mode

Thread

Thread of execution:

an abstraction that contains enough state of a running program that it can be stopped and resumed

xv6 API: yield, swtch

Goals for solution:

Switching transparent to user threads

User thread cannot hog a processor (kernel thread assumed to be correct, so not a goal)

Overview of switch between two user threads

user threads

User -> kernel transition

kernel -> kernel switch

kernel -> User transition

guaranteed U->K transitions

timing interrupt every 100 ms

switches to different kernel thread on yield

the different kernel thread returns to a different user thread

Challenges in implementing:

Opaque code ("You are not supposed to understand this")

Concurrency (several processors switching between threads)

Terminating a thread, always need a valid stack

Xv6 design

One scheduler thread per processor

Simplifies implementation:

Need a stack to run scheduler on, if there are no other threads anymore

Downside: more switches

To switch from one thread to another requires two switches

thread 1 -> scheduler -> thread 2

Advantage: scheduling organized as co-routines

Simplifies reasoning about concurrency

E.g., it is clear lock is always passed from current thread to scheduler thread

xv6 schedules user threads preemptively

Every 100ms a timer interrupt

xv6 schedules kernel threads cooperatively

Assume kernel programmer doesn't make mistakes (e.g., no infinite loop)

Code

Forced switching:

demo of two processes who don't invoke system calls

look at process states

proc.h

clock interrupt

lapic.c for SMP

timer.c for uniprocessor

walk through what xv6 does to guarantee switching

breakpoint in trap at yield() (b: trap.c:136)

get hog running (c 100)

look at tf, in particular tf->eip

interrupt arrived in user space, while running hog

look at tf->trapno (timer interrupt), gets to yield

get to swtch, look at contexts (p /x *cpus[0]->scheduler)

Q: why should ncli be 1?

look at eip before return from swtch (we switched to scheduler thread)

scheduler: switches to selected thread (set b proc.c:288)

will return user space

What is the scheduling policy?

will the thread that called yield run immediately again?

Concurrency

plock held across swtch; why?

yield(): p is set to runnable, p must complete switch before another scheduler choses p
hard to reason about; coroutine style helps

always enter and leave a thread at the same instruction

can two schedulers select the same runnable process?

why does scheduler release after loop, and re-acquire it immediately?

- To give other processors a chance to look at the proc table

why does the scheduler enable interrupts?

- The scheduler maybe the only runnable thread

- All other threads maybe sleeping (e.g., waiting on I/O input)

- To briefly run the processor with interrupts enabled

- An interrupt may wakeup a sleeping thread

Thread clean up

let's look at kill()

- can we clean up killed process? (no: it might be running, holding locks, etc.)

- before returning to user space: process kills itself by calling exit

let's look at exit()

- can thread delete its stack?

- no: it has to switch off it!

- wait() does the cleanup