

6.828 2014 L15: Operating System Organization

=

Operating System Organization

students have completed building an exokernel in lab3 and lab4
 user-space fork (copy-on-write)
 sophisticated VM handling in libos
 assumption in this lecture: one cannot change the kernel

Plan: OS organization

goals for a kernel interface
 monolithic
 microkernel
 exokernel
 little data, much opinion

OS organization

--

Goals for kernel

Apps can use hardware resources
 Apps are easy convenient to write
 Apps are multiplexed (isolation and sharing)
 Few kernel crashes
 Kernel can evolve

OS design

lots of ways to structure an OS -- how to decide?
 what is the right kernel API?
 assumption: you cannot change the API radically
 => the point of a view in this lec: the app developer
 looking for principles and approaches

Traditional approach (Linux, xv6)

kernel API ~ POSIX
 virtualize some resources: cpu and memory
 simulate a dedicated cpu and memory system for each app
 why? it's a simple model for app programmers
 abstract others: storage, network, IPC
 layer a sharable abstraction over h/w (file system, IP/TCP)

Example: virtualize the cpu

goal: simulate a dedicated cpu for each process
 we want transparent CPU multiplexing
 process need not think about how it interacts w/ other processes
 OS runs different processes in turn, via clock interrupt
 clock means process doesn't need to do anything special to switch
 also prevents hogging
 how to achieve transparency?
 OS saves state, then restores
 what does OS save?
 eight regs, EIP, seg regs, eflags, page table base ptr
 where does OS save it?
 OS keeps per-process table of saved states
 the return from clock interrupt restores a *different* process's state
 the point: process doesn't have to worry about multiplexing!
 this is the traditioal approach to virtualizing the CPU
 what does the exokernel/JOS do?

Example: virtualize memory

idea: simulate a complete memory system for each process
 so process has complete freedom how it uses that memory
 doesn't have to worry about other processes
 so addresses $0..2^{32}$ all work, but refer to private memory

convenient: all programs can start at zero
 and memory looks contiguous, good for large arrays &c
 safe: can't even *name* another process's memory
 again: traditional but we'll soon see it's a limiting approach
 really want apps to have more control than this style of VM implies

Level of indirection allows OS to play other tricks

demand paging:

process bigger than available physical memory?
 "page-out" (write) pages to disk, mark PTEs invalid
 if process tries to use one of those pages, MMU causes page fault
 kern finds phys mem, page-in from disk, mark PTE valid
 this works because apps use only a fraction of mem at a given time
 need h/w valid flag, page faults, and re-startable instructions

copy-on-write:

avoid copy implied by fork() -- won't be needed if exec()
 make parent and child share the physical memory pages
 if either writes, do the copy then
 so need per-page write-protect flag
 both of above are transparent to application
 still thinks it has simple dedicated memory from 0..2³²

paging h/w has turned out to be one of the most fruitful ideas in OS
 you have been using it a lot in labs
 can we make it safe for apps to play these tricks?

Monolithic

--

Traditional organization: monolithic OS (e.g., xv6)

h/w, kernel, user

kernel is a big program: process ctl, vm, fs, network
 all of kernel runs w/ full hardware privilege (very convenient)
 good: easy for sub-systems to cooperate (e.g. paging and file system)
 bad: interactions => complex, bugs are easy, no isolation within OS
 extensibility: dynamically-loadable modules, wait for next kernel release
 philosophy: convenience (for app or OS programmer)
 for any problem, either hide it from app, or add a new system call
 (we need philosophy because there is not much science here)
 may take a while for the new system call is added
 very successful approach (e.g., Linux)

Microkernel

--

Alternate organization: microkernel

philosophy: IPC and user-space servers

for any problem, make a new server, talk to it w/ RPC

h/w, kernel, server processes, apps

servers: VM, FS, TCP/IP, Print, Display

split up kernel sub-systems into server processes

some servers have privileged access to some h/w (e.g. FS and disks)

apps talk to them via IPC / RPC

kernel's main job: fast IPC

good: simple/efficient kernel, sub-systems isolated, enforced better modularity

bad: cross-sub-system optimization harder, lots of IPCs may be slow

extensibility: can replace servers

in the end, lots of good individual ideas, but overall plan didn't catch on for desktops/servers

but success in embedded space

More next lecture

Exokernel

--

Alternate organization: exokernel (JOS)

philosophy: eliminate all abstractions

for any problem, expose h/w or info to app, let app do what it wants

h/w, kernel, environments, libOS, app

an exokernel would not provide address space, virtual cpu, file system, TCP

instead, give control to app:

phys pages, addr mappings, clock interrupts, disk i/o, net i/o

let app build nice address space if it wants, or not

should give aggressive apps much more flexibility

challenges:

how to multiplex cpu/mem/&c if you expose directly to apps?

how to get security/isolation despite apps having low-level control?

how to multiplex w/o understanding: disk (file system), incoming tcp pkts

Exokernel example: memory

what are the resources? (phys pages, mappings)

what does an app need to ask the kernel to do?

pa = AllocPage()

DeallocPage(pa)

TLBwr(va, pa)

and these kernel->app upcalls:

PageFault(va)

PleaseReleaseAPage()

what does OS need to do to make multiplexing work?

ensure app only creates mappings to phys pages it owns

track what env owns what phys pages

decide which app to ask to give up a phys page when system runs out

that app gets to decide which of its pages

Simple example: shared memory

two processes want to share memory, for fast interaction

note traditional "virtual address space" doesn't allow for this

process a: pa = AllocPage()

put 0x5000 -> pa in private table

PageFault(0x5000) upcall -> TLBwr(0x5000, pa)

give pa to process b (need to tell exokernel...)

process b:

put 0x6000 -> pa in private table

...

Example cool thing you could do w/ exokernel-style memory

databases like to keep a cache of disk pages in memory

problem on traditional OS:

assume an OS with demand-paging to/from disk

if DB caches some disk data, and OS needs a phys page,

OS may page-out a DB page holding a cached disk block

but that's a waste of time: if DB knew, it could release phys

page w/o writing, and later read it back from DB file (not paging area)

1. exokernel needs phys mem for some other app

2. exokernel sends DB a PleaseReleaseAPage() upcall

3. DB picks a clean page, calls DeallocPage(pa)

4. OR DB picks dirty page, writes to disk, then DeallocPage(pa)

Exokernel example: cpu

what does it mean to expose cpu to app?

kernel tells app when it is taking away cpu

kernel tells app when it gives cpu to app

so if app is running and timer interrupt causes end of slice

cpu jumps from app into kernel

kernel jumps back into app at "please yield" upcall

app saves state (registers, EIP, &c)

app calls Yield()

when kernel decides to resume app

kernel jumps into app at "resume" upcall

app restores those saved registers and EIP

What cool things could an app do w/ exokernel-style cpu management?
suppose time slice ends in the middle of
 acquire(lock);
 ...
 release(lock);
you don't want the app to be holding the lock the whole time!
then maybe other apps can't make forward progress
so the "please yield" upcall can first complete the critical section

Fast RPC with direct cpu management

how does traditional OS let apps communicate?
pipes (or sockets)
picture: buffer in kernel, lots of copying and system calls
RPC probably takes 8 kernel/user crossings (read()s and write()s)
how does exokernel help?
Yield() can take a target process argument
 almost a direct jump to an instruction in target process
 kernel allows only entries at approved locations in target
 kernel leaves regs alone, so can contain arguments
 (in contrast to traditional restore of target's registers)
target app uses Yield() to return
so only 4 crossings

Debate: monolithic versus microkernel

Get OS developers very excited

Practical view:

 If you have a working monolithic kernel, why change?
 Adding features that users want
 Isolate new features, if possible
 If you start from scratch, why not use microkernel?

Today: systems are not purely one or the other design

Linux is monolithic but has servers
Linux API also gives apps a lot of control
 map/unmap, pgfault
 kernel extensions